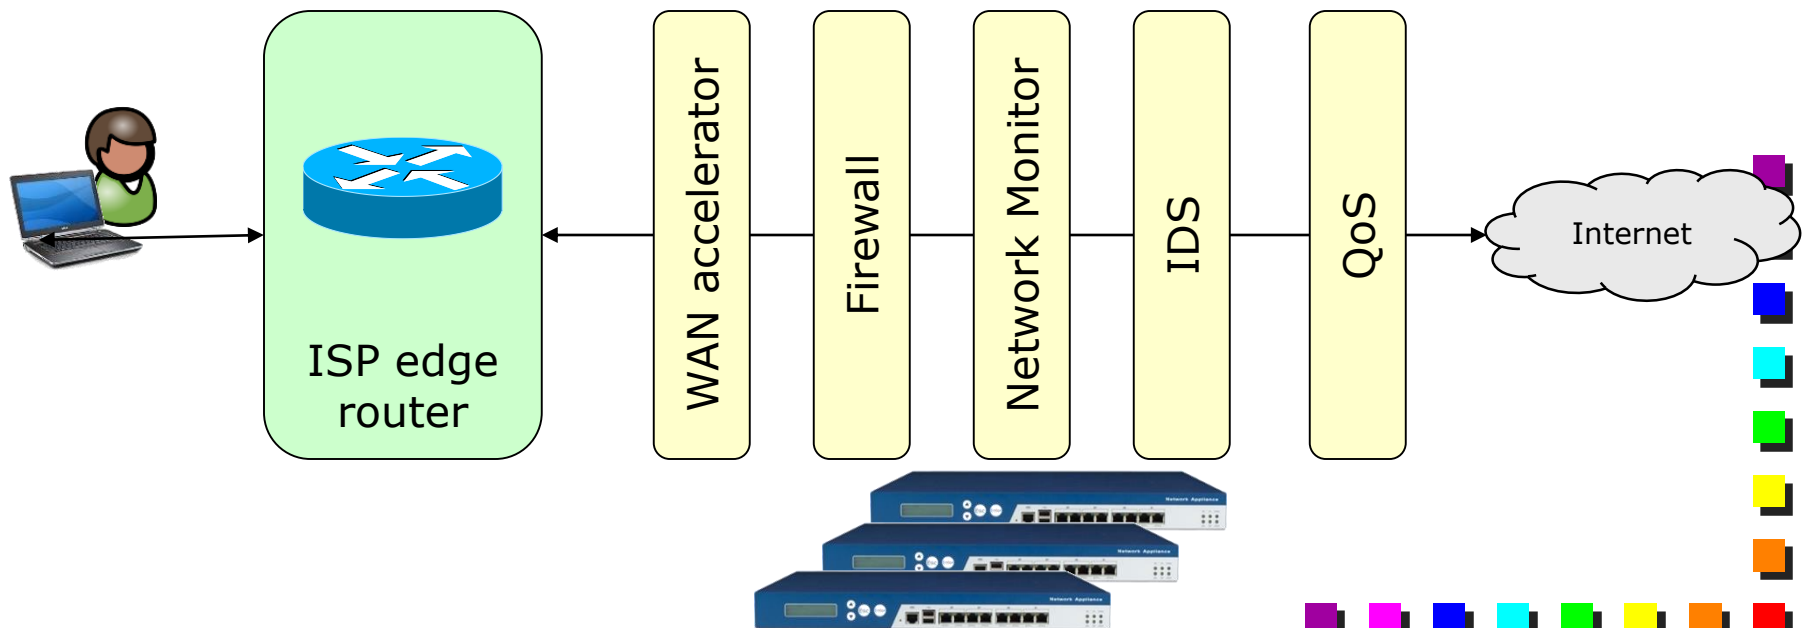


Network Functions Virtualization

Fulvio Riso
Politecnico di Torino

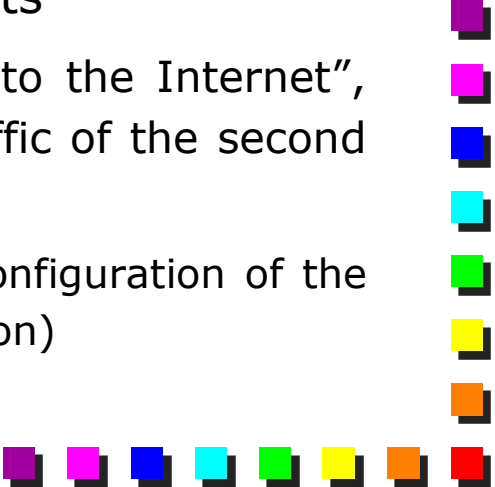
Service Function Chaining

- Often, particularly at the edge of the network, we need to chain **different dedicated hardware appliances** to provide added-value services
- This is what is called a **chain of network functions**

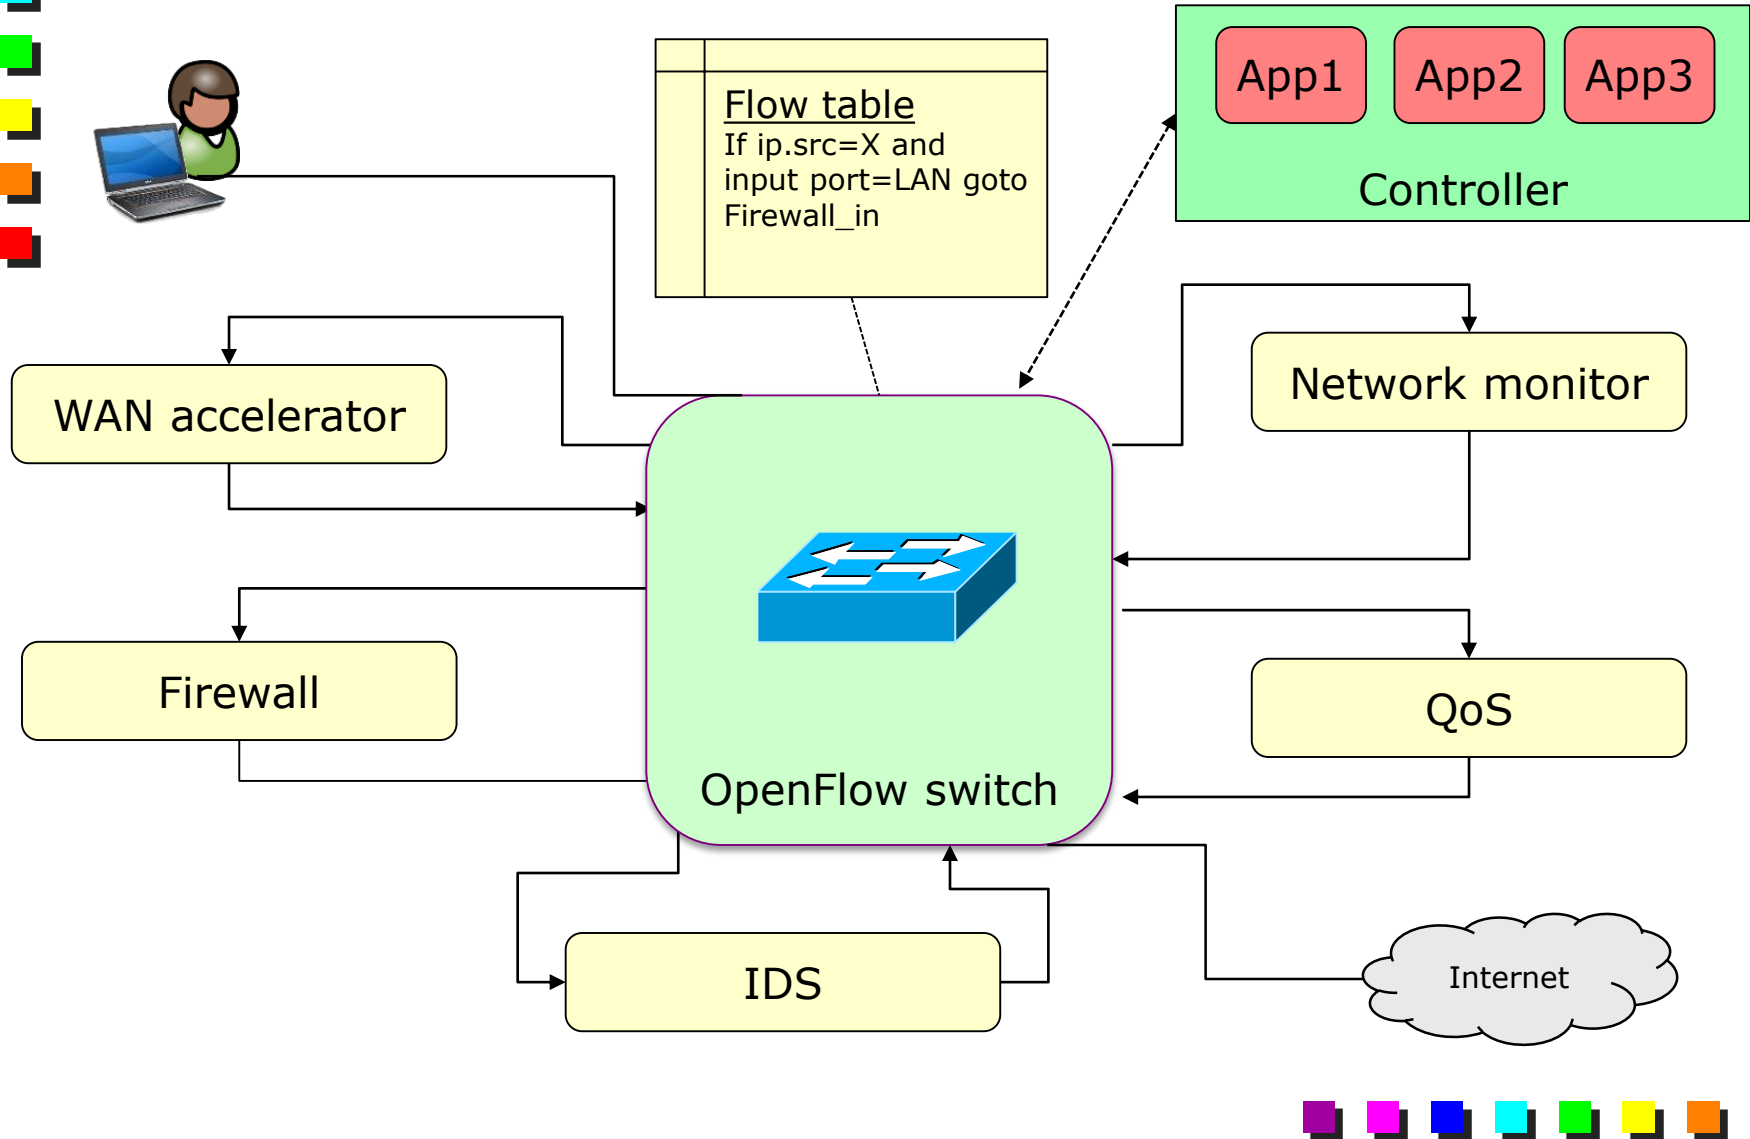




Several (practical) problems with SFC

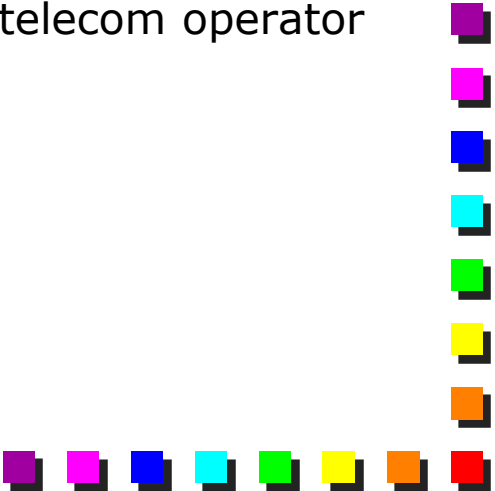
- Hardware resources not used at best
 - Some appliances may sustain an heavy load, while other may be almost unloaded and we are not able to share the available hardware resources (e.g., CPU, memory) between different services
 - Service disruption when modifying the service chain
 - Each time we add/remove a middlebox, we have to disrupt the service
 - Not easy to differentiate services among tenants
 - What about it a tenant buys a “secure access to the Internet”, but other don’t? How can we avoid that the traffic of the second tenant goes through the firewall as well?
 - This requires the firewall to support explicit configuration of the user privileges (i.e., per-application configuration)
- 

Service Function Chaining with SDN (1)



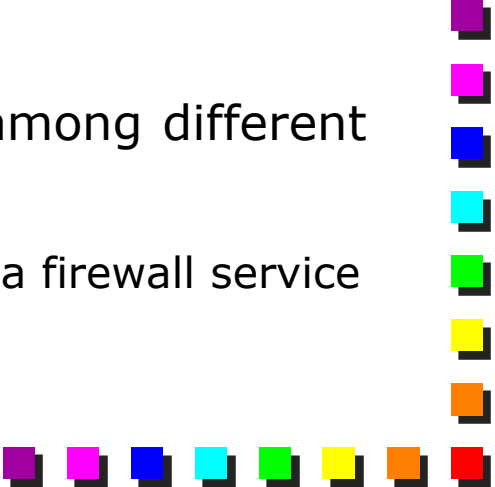


Service Function Chaining with SDN (2)

- An OpenFlow switch can be installed to connect all boxes together
 - OpenFlow rules can be used to steer the traffic from each user to the proper set of services
 - Rules can be either pre-provisioned, or provisioned on demand (e.g., user logs-in, and the controller instantiates the proper rules for this user, valid only for the duration of the user session)
 - The controller can be installed locally to the machines
 - This looks like a nice setup for an edge POP of a telecom operator
- 



SFC with SDN: characteristics

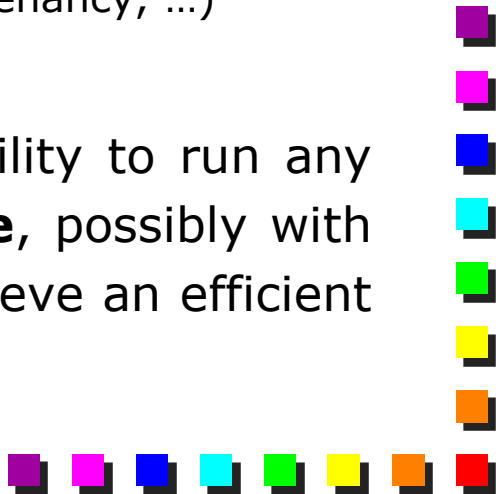
- Agility in provisioning new services
 - Install the box, then “routing” is done via software instead of connecting the box to the other with physical wires
 - Maintenance and reliability
 - Cabling is done once
 - Different customers can have different service chains
 - “routing” done via software, even possible to change its decisions based on other parameters (e.g., application layer content)
 - Still difficult to partition a physical appliance among different tenants
 - Many small business customers, each asking for a firewall service
- 



Network Functions Virtualization (1)

- Four main components:

- Fast standard hardware (e.g. Intel servers)
 - Commercial-off-the-shelf (COTS) hardware
- Software-based network functions
 - Network functions, previously running on a dedicated appliance, now become a software image, running on a standard server
- Computing virtualization (e.g., Linux KVM)
 - All advantages of virtualization (quick provisioning, scalability, mobility, reduced CapEx, reduced OpEx, multitenancy, ...)
- Standard API (i.e., ETSI framework)

- Network Functions Virtualization is the capability to run any **network function** on a **standard hardware**, possibly with the help of **computing virtualization** to achieve an efficient use of resources
- 



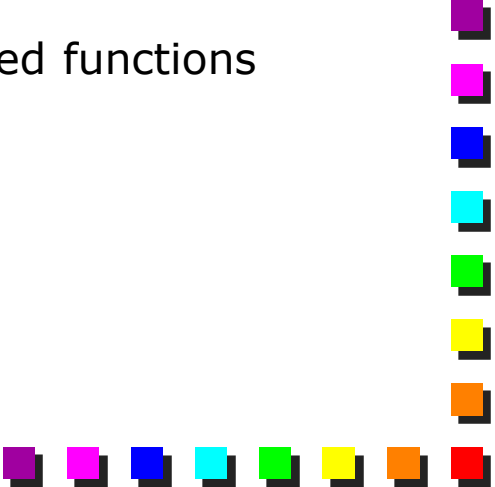
Network Functions Virtualization (2)

- Two possible deployment scenario for NFV services

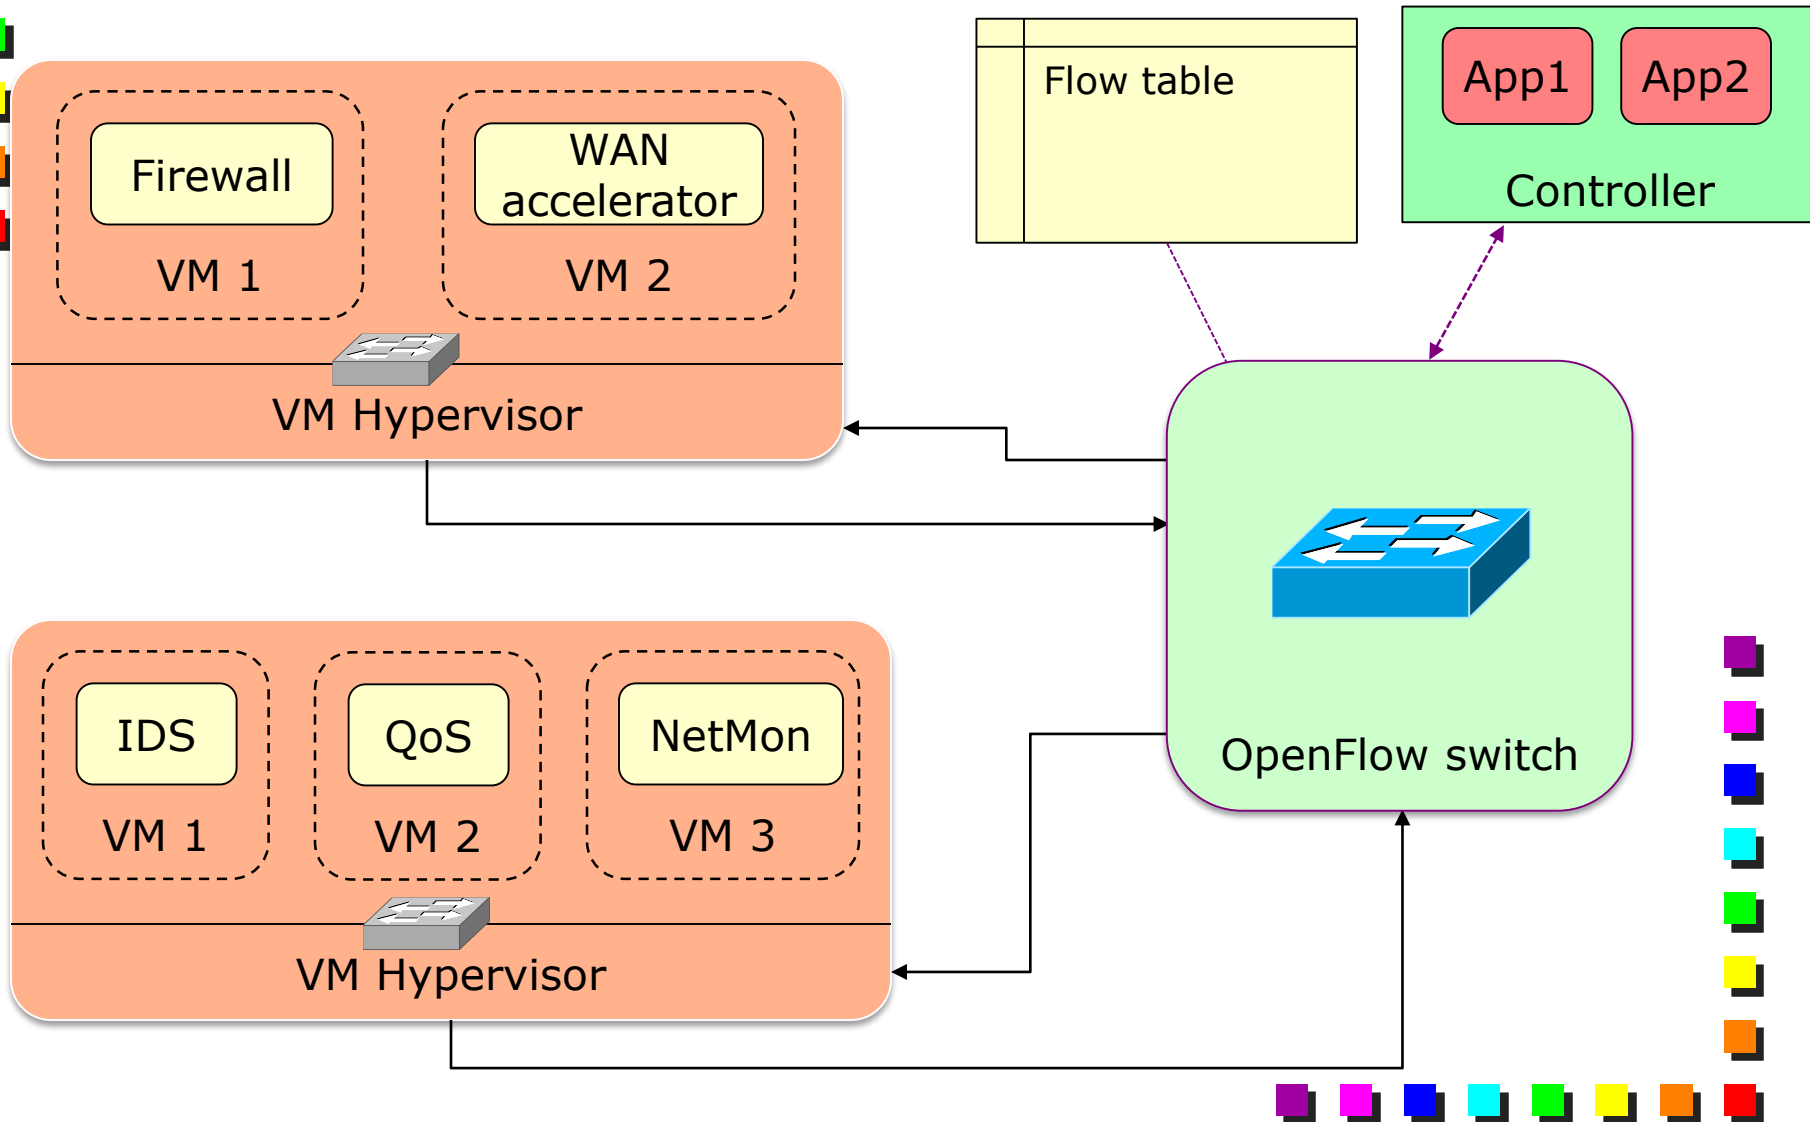
- **Software based Devices**

- Instead of having the service in a dedicated appliance, the service runs on standard hardware
 - E.g., Routers, Firewalls, Broadband Network Gateways (BNG) in a *white box* implementation
- Often, virtualization is not used in this case (or used internally, without allowing the server to be integrated in the datacenter of the provider)
- Commonly created through the use of DPDK-based functions

- **Function Modules**


- Refers to both data plane and control plane
 - E.g., DHCP, NAT, Rate Limiting, etc.
 - Often they come as pure software packages
- 

Service functions chaining with NFV






Advantages of NFV


- **1. Virtualization:** use resources without worrying about where it is physically located, how much it is, how it is organized, etc.
 - **2. Orchestration:** manage thousands of devices
 - **3. Programmable:** can change the behavior on the fly
 - **4. Dynamic Scaling:** can adapt to different workloads
 - **5. Automation**
 - **6. Visibility:** Monitor resources, connectivity
 - **7. Performance:** Optimize network device utilization
 - **8. Multi-tenancy**
 - **9. Service Integration**
 - **10. Openness:** Full choice of service modules
- 

Partially adapted from http://www.cse.wustl.edu/~jain/cse570-13/m_17nfv.htm



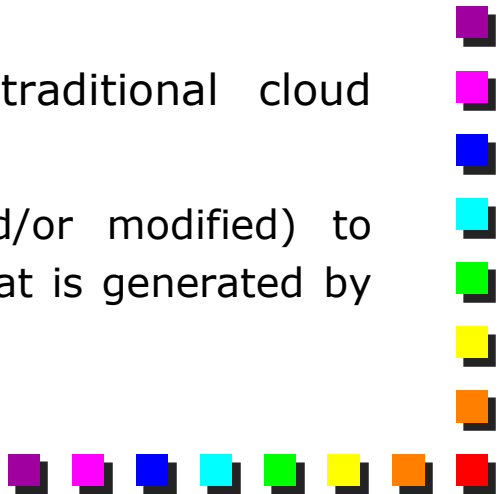


Chaining vs general services (in NFV)

- Chaining usually refers to a service that is made up of a **stack** of modules
 - Services are not always stackable
 - E.g., DHCP, DNS, web services need to operate in a LAN
 - How to model a LAN with a chain?
 - Hence, NFV needs to be more flexible than just support chains
- 





NFV and cloud

- NFV can be seen as a way to bring network services in the world of cloud technologies
 - Cloud: hosts web servers, database servers, big data applications, etc.
 - NFV: adds also network services to that picture
 - Although apparently NFV can be realized mostly with existing technologies, in practice:
 - Cloud frameworks may not support well traffic steering, although they support well traditional LAN services
 - Network services are I/O intensive, while traditional cloud services are mostly CPU intensive
 - Some technologies need to be tuned (and/or modified) to support the high amount of network traffic that is generated by network services
- 

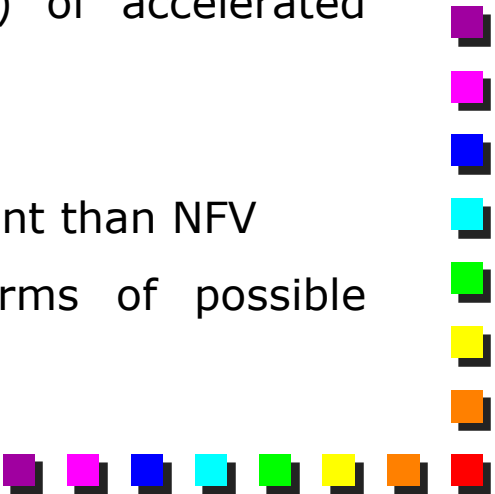


NFV and SDN (1)

- NFV is about computing, SDN is about network paths
 - NFV requires SDN for flexible traffic steering
 - Although, a point-to-point Ethernet is often enough for most of the purposes
 - NFV and SDN are complementary
 - One does not depend upon the other
 - You can do SDN only, NFV only, or SDN and NFV
 - A lot of discussions about SDN, not much debate about NFV
- 
- 




NFV and SDN (2)

- Both have similar goals but approaches are very different
 - SDN needs new interfaces, control modules, applications must re-engineered
 - NFV requires moving network applications from dedicated hardware to virtual images on standard hardware
 - SDN heavily leverages accelerated hardware (the hardware switch)
 - NFV can hardly take advantages (right now) of accelerated hardware
 - Hence, SDN can be potentially much more efficient than NFV
 - NFV is currently much more flexible (in terms of possible supported applications) than SDN
- 


VNF and network traffic (1)

- In theory, VMs can be deployed based on the resources that are available on the data center
- In practice, this may lead to very un-optimized paths

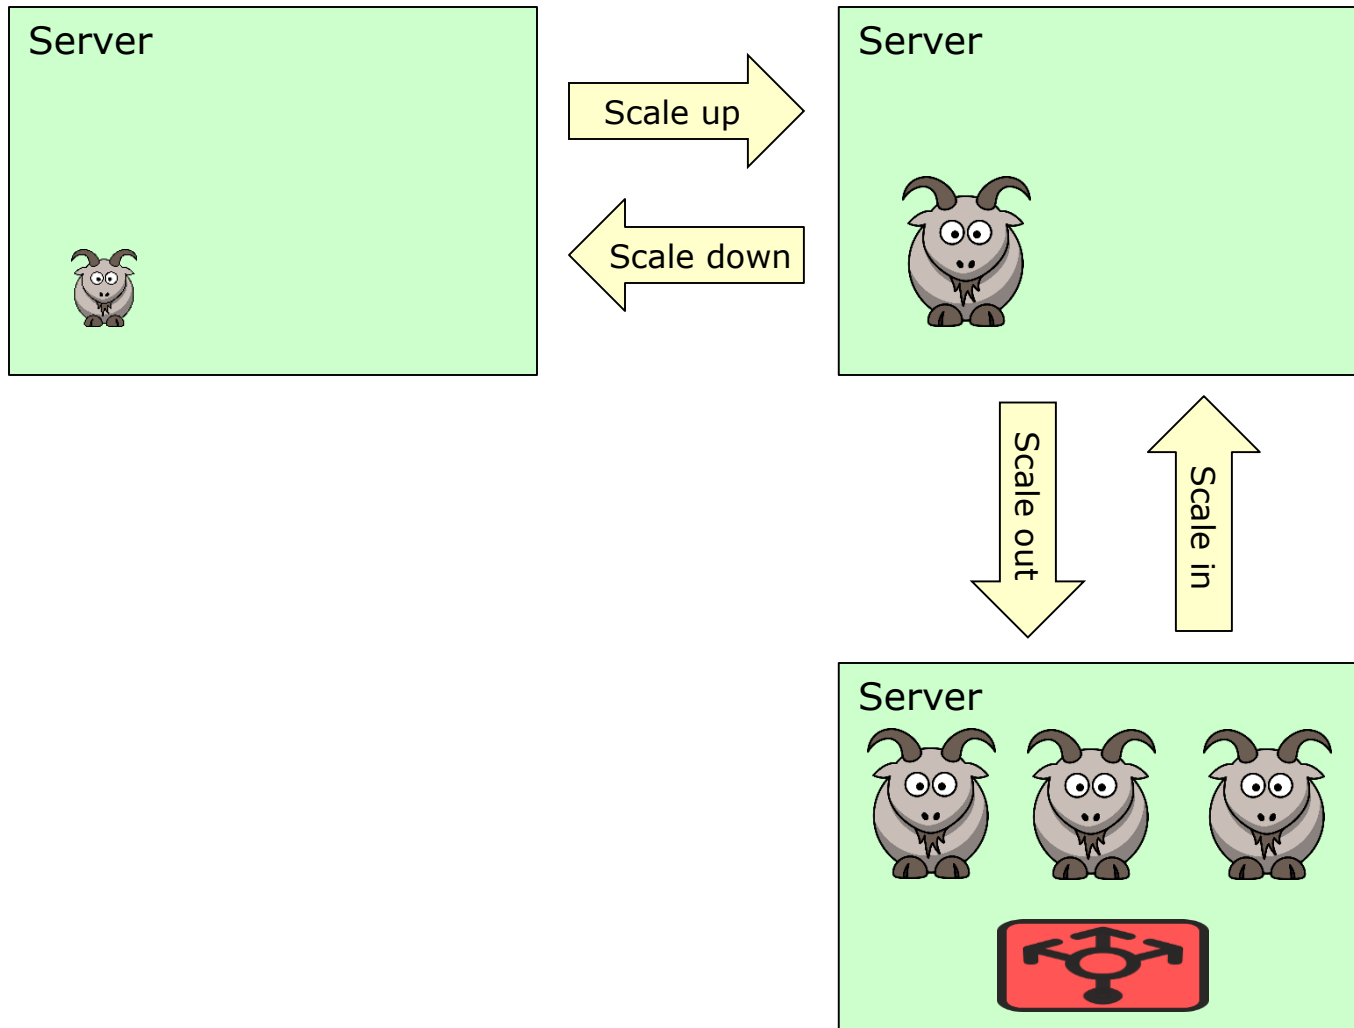




VNF and network traffic (2)

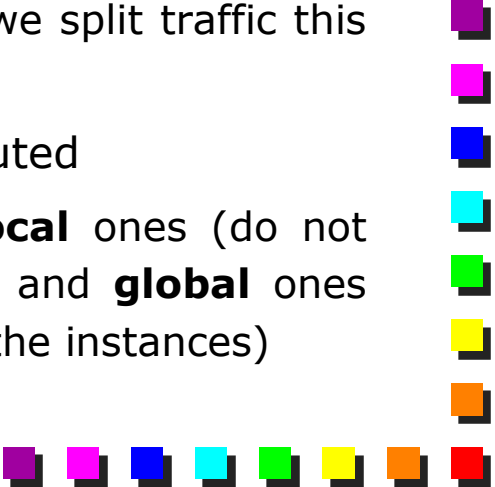
- NFV may have a huge impact on the traffic of your datacenter
 - NFV can generate a huge amount of traffic on the **network**
 - NFV can generate a huge amount of traffic **inside each server** as well
 - Packets may travel several times back and forth to the switch
 - We may need to optimize the computing technologies to reduce the load (e.g., SR-IOV, VirtIO, Shared memory)
 - We need to predict the amount of traffic that is generated in the datacenter to avoid troubles
- 

Scaling: definitions





NFV and scalability

- VMs are good when we need to **consolidate** many (tiny) application instances on the same physical servers
 - VMs are not very good when an application requires so many resources that even a **fully dedicated server is not enough** to deliver the service
 - In the latter case, we have mainly two options:
 - Add a load balancer (e.g., using SDN) in front of the different instances and make sure that they can operate independently
 - Most applications work per-TCP-session, so if we split traffic this way, the application can operate properly
 - Modify the application in order to make it distributed
 - Application can have two set of variables: **local** ones (do not need to be in sync with the other instances) and **global** ones (each modification has to be propagated to all the instances)
- 



The ETSI NFV model

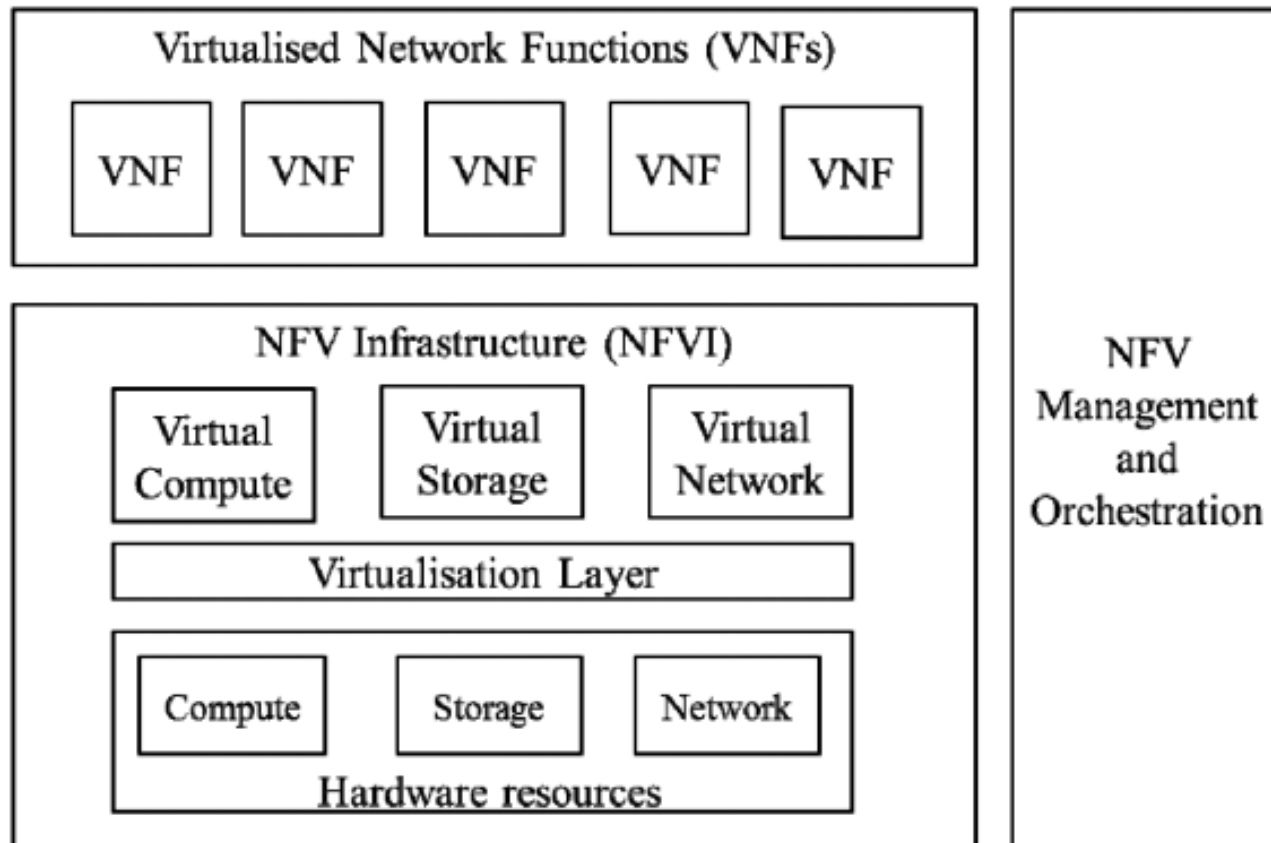




ETSI NFV ISG


- ETSI NFV Industry Specification Group (ISG): define the requirements, architectural framework, interfaces for NFV
 - <http://www.etsi.org/technologies-clusters/technologies/nfv>
 - Different Working Groups / Expert groups
 - Architecture for the virtualization infrastructure
 - Management and orchestration
 - Software architecture
 - Reliability and availability, resilience and fault tolerance
 - Public demonstrations and Proof of Concept
 - Performance
 - Security
- 

High-level NFV framework





NFV terminology (1)


- **Network Function (NF):** functional building block with a well defined interfaces and well defined functional behavior
 - **Virtualized Network Function (VNF):** software implementation of NF that can be deployed in a virtualized infrastructure
 - **VNF Set:** connectivity between VNFs is not specified, e.g., residential gateways
 - **VNF Forwarding Graph:** service chain when network connectivity order is important, e.g., firewall, NAT, load balancer
 - **NFV Infrastructure (NFVI):** hardware and software required to deploy, manage and execute VNFs including computation, networking, and storage
- 



NFV terminology (2)

- **NFVI Point of Presence (PoP):** location of NFVI
- **NFVI-PoP Network:** internal network
- **Transport Network:** network connecting a PoP to other PoPs or external networks
- **VNF Manager:** VNF lifecycle management e.g., instantiation, update, scaling, query, monitoring, fault diagnosis, healing, termination
- **Virtualized Infrastructure Manager:** management of computing, storage, network, software resources
- **Network Service:** a composition of network functions and defined by its functional and behavioral specification
- **NFV Service:** a network services using NFs with at least one VNF

Partially adapted from http://www.cse.wustl.edu/~jain/cse570-13/m_17nfv.htm



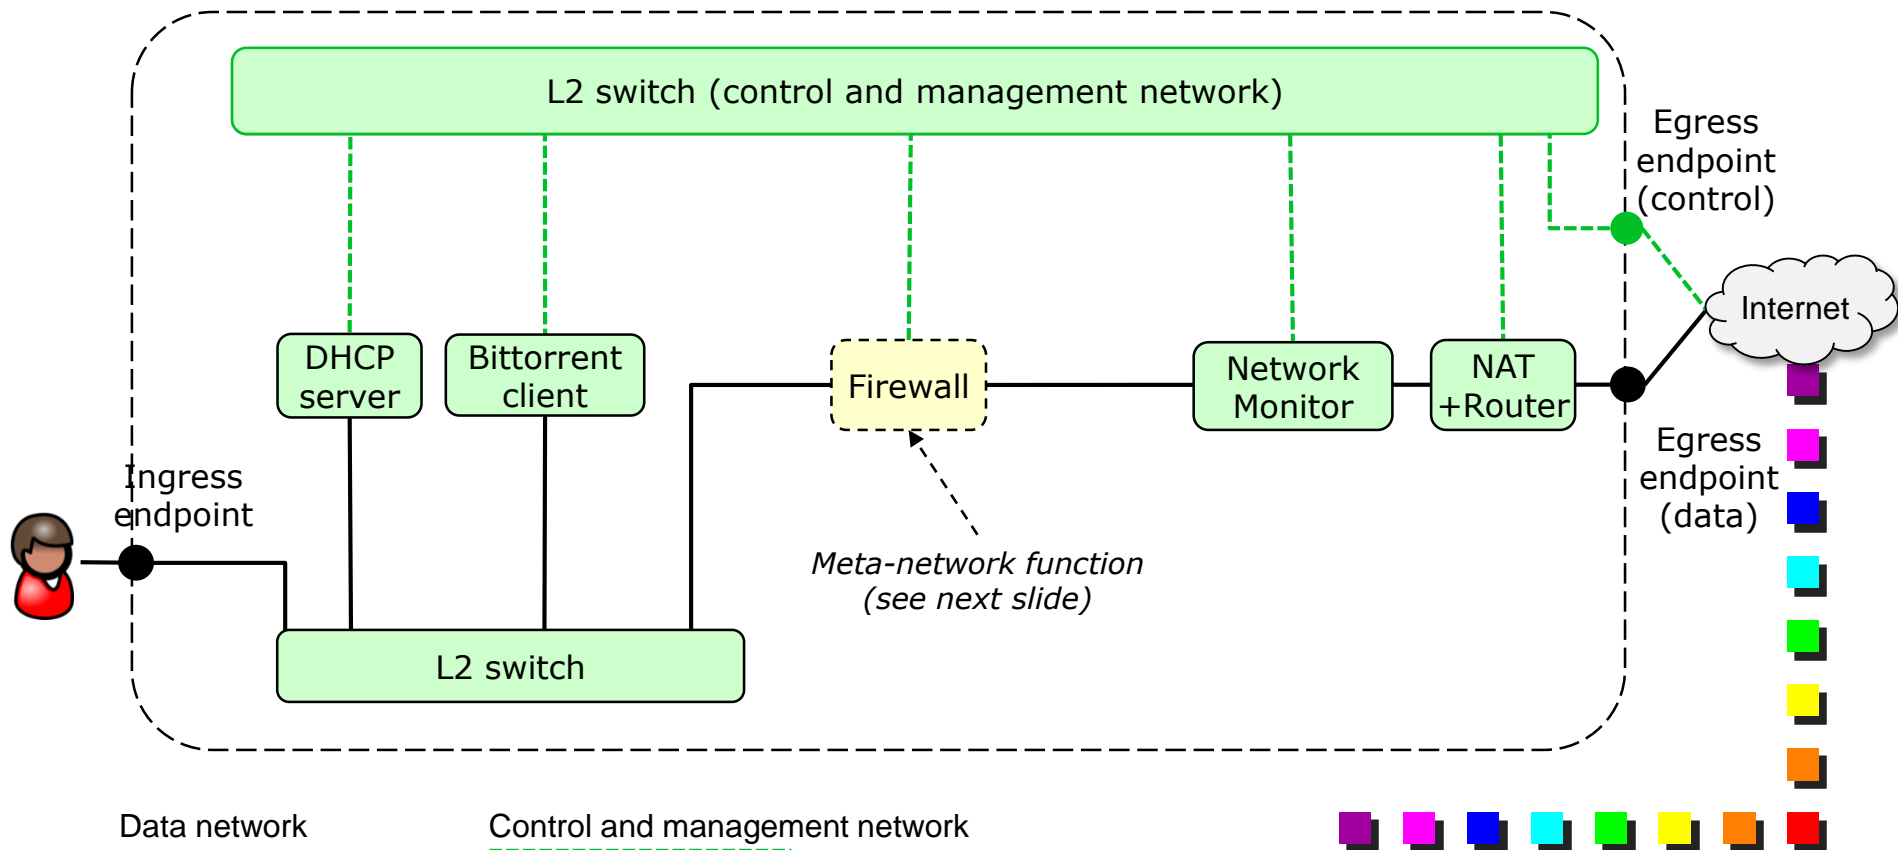


NFV terminology (3)

- **User Service:** services offered to end users / customers / subscribers
- **Deployment Behavior:** NFVI resources required by a VNF, e.g., number of VMs, memory, disk, images, bandwidth, latency
- **Operational Behavior:** VNF instance topology and lifecycle operations, e.g., start, stop, pause, migration, ...
- **VNF Descriptor:** deployment behavior + operational behavior
- **NFV Orchestrator:** automates the deployment, operation, management, coordination of VNFs and NFVI
- **VNF Forwarding Graph:** connection topology of various NFs of which at least one is a VNF

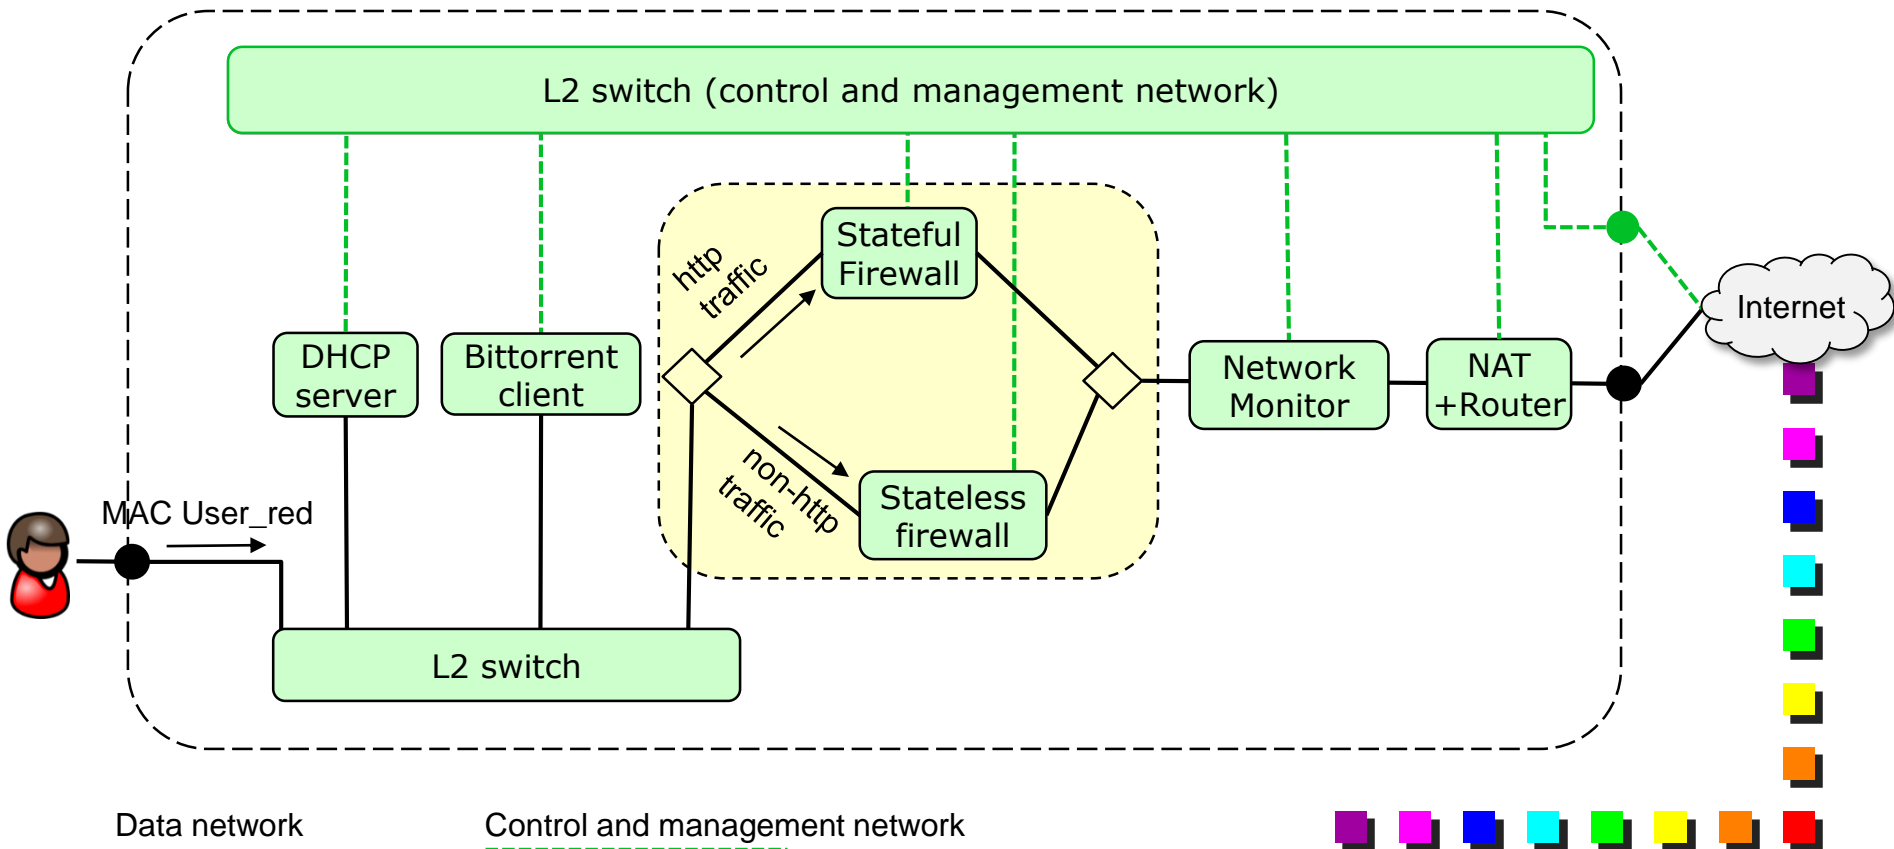
Forwarding graph

- High-level representation of the service in terms of functional blocks and their connections, similar to a service chain
- Example of a complex service

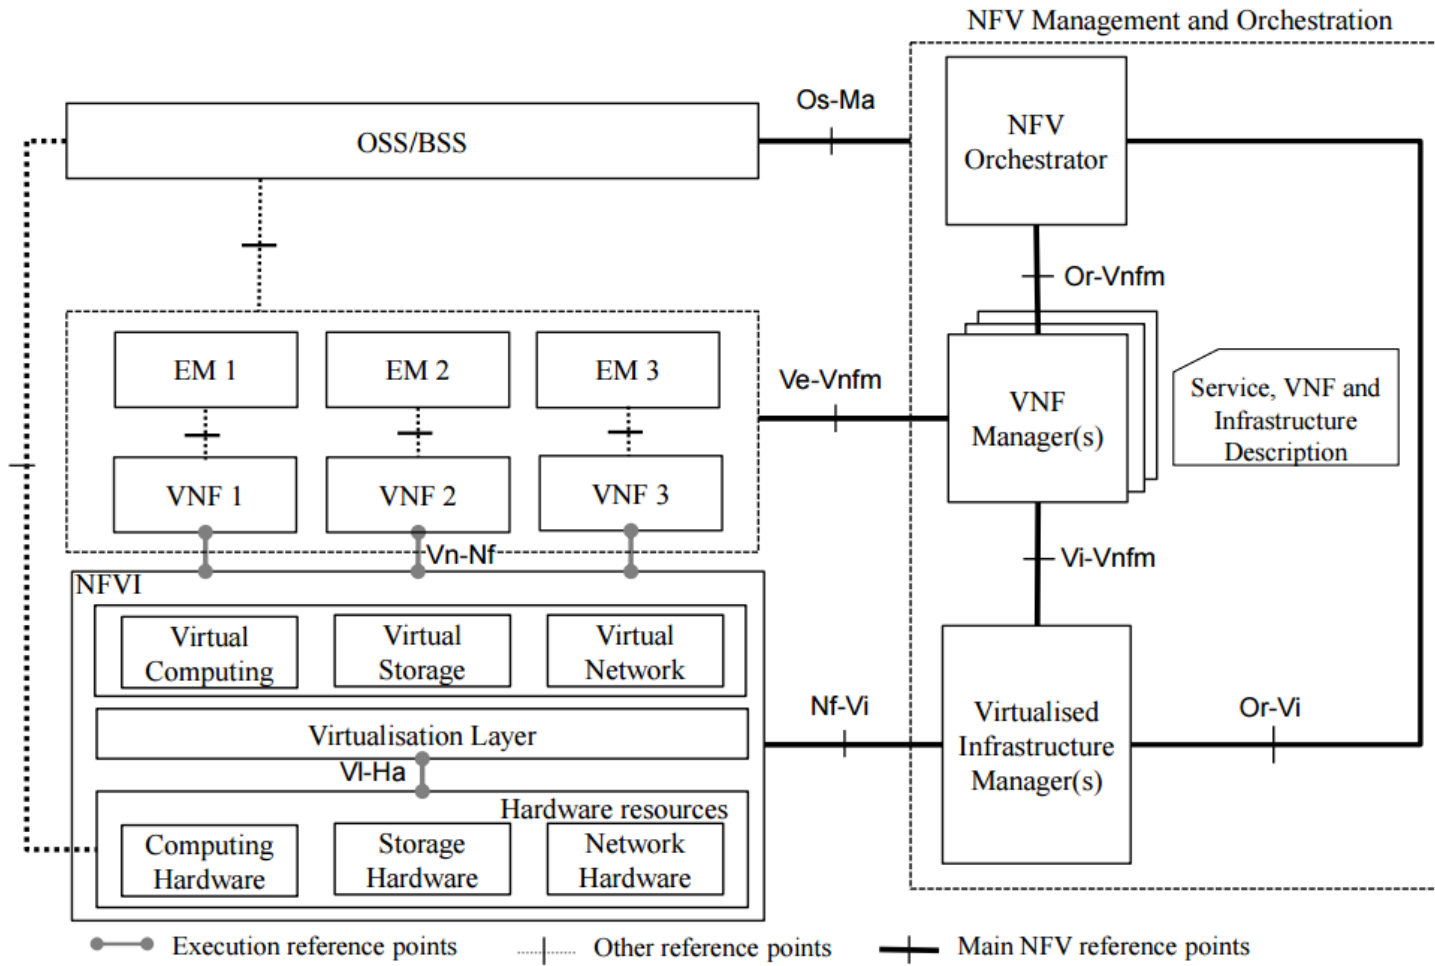


Forwarding graph: hierarchical decomposition

- Services can be hierarchically decomposed in smaller building blocks



NFV Reference Architectural Framework

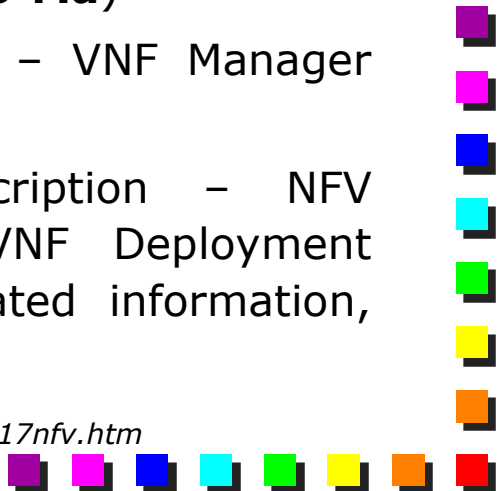




NFV Reference Points

- Reference Point: points for inter-module specification
 - 1. Virtualization Layer-Hardware Resources (**VI-Ha**)
 - 2. VNF – NFVI (**Vn-Nf**)
 - 3. Orchestrator – VNF Manager (**Or-Vnfm**)
 - 4. Virtualized Infrastructure Manager – VNF Manager (**Vi-Vnfm**)
 - 5. Orchestrator – Virtualized Infrastructure Manager (**Or-Vi**)
 - 6. NFVI-Virtualized Infrastructure Manager (**Nf-Vi**)
 - 7. Operation Support System (OSS)/Business Support Systems (BSS) – NFV Management and Orchestration (**Os-Ma**)
 - 8. VNF/ Element Management System (EMS) – VNF Manager (**Ve-Vnfm**)
 - 9. Service, VNF and Infrastructure Description – NFV Management and Orchestration (**Se-Ma**): VNF Deployment template, VNF Forwarding Graph, service-related information, NFV infrastructure information

Partially adapted from http://www.cse.wustl.edu/~jain/cse570-13/m_17nfv.htm






NFV Framework Requirements

- **1. General:** partial or full virtualization, predictable performance
- **2. Portability:** decoupled from underlying infrastructure
- **3. Performance:** as described and facilities to monitor
- **4. Elasticity:** scalable to meet SLAs; movable to other servers
- **5. Resiliency:** be able to recreate after failure; specified packet loss rate, calls drops, time to recover, etc.
- **6. Security:** role-based authorization, authentication
- **7. Service Continuity:** seamless or non-seamless continuity after failures or migration
- **8. Service Assurance:** time stamp and forward copies of packets for fault detection
- **9. Energy Efficiency Requirements:** should be possible to put a subset of VNF in a power conserving sleep state
- **10. Transition:** coexistence with legacy and interoperability among multi-vendor implementations
- **11. Service Models:** operators may use NFV infrastructure operated by other operators

Partially adapted from http://www.cse.wustl.edu/~jain/cse570-13/m_17nfv.htm






An example of a complex service using NFV



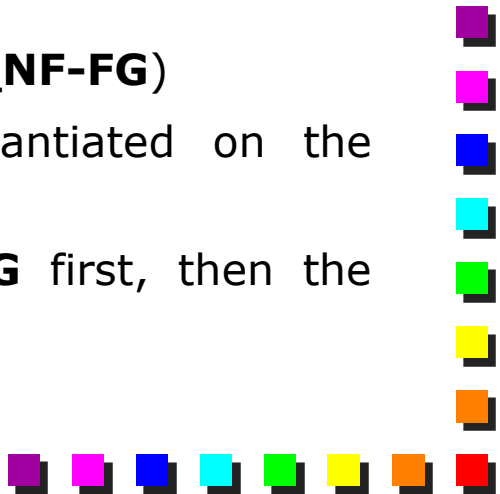
Service overview

- Different users experiment a different and customized network service based on their credentials
 - User **RED** deploys a **RED_NF-FG** that includes a **RED_SFTP** server and a **RED_FIREWALL**
 - User **BLUE** deploys a **BLUE_NF-FG** that includes a **BLUE_SFTP** server and a **BLUE_FIREWALL**
- Network operator sets up some additional services in the network, active on all users




A vertical column of colored squares (purple, magenta, blue, cyan, green, yellow, orange, red) on the left side of the slide.

Demo details

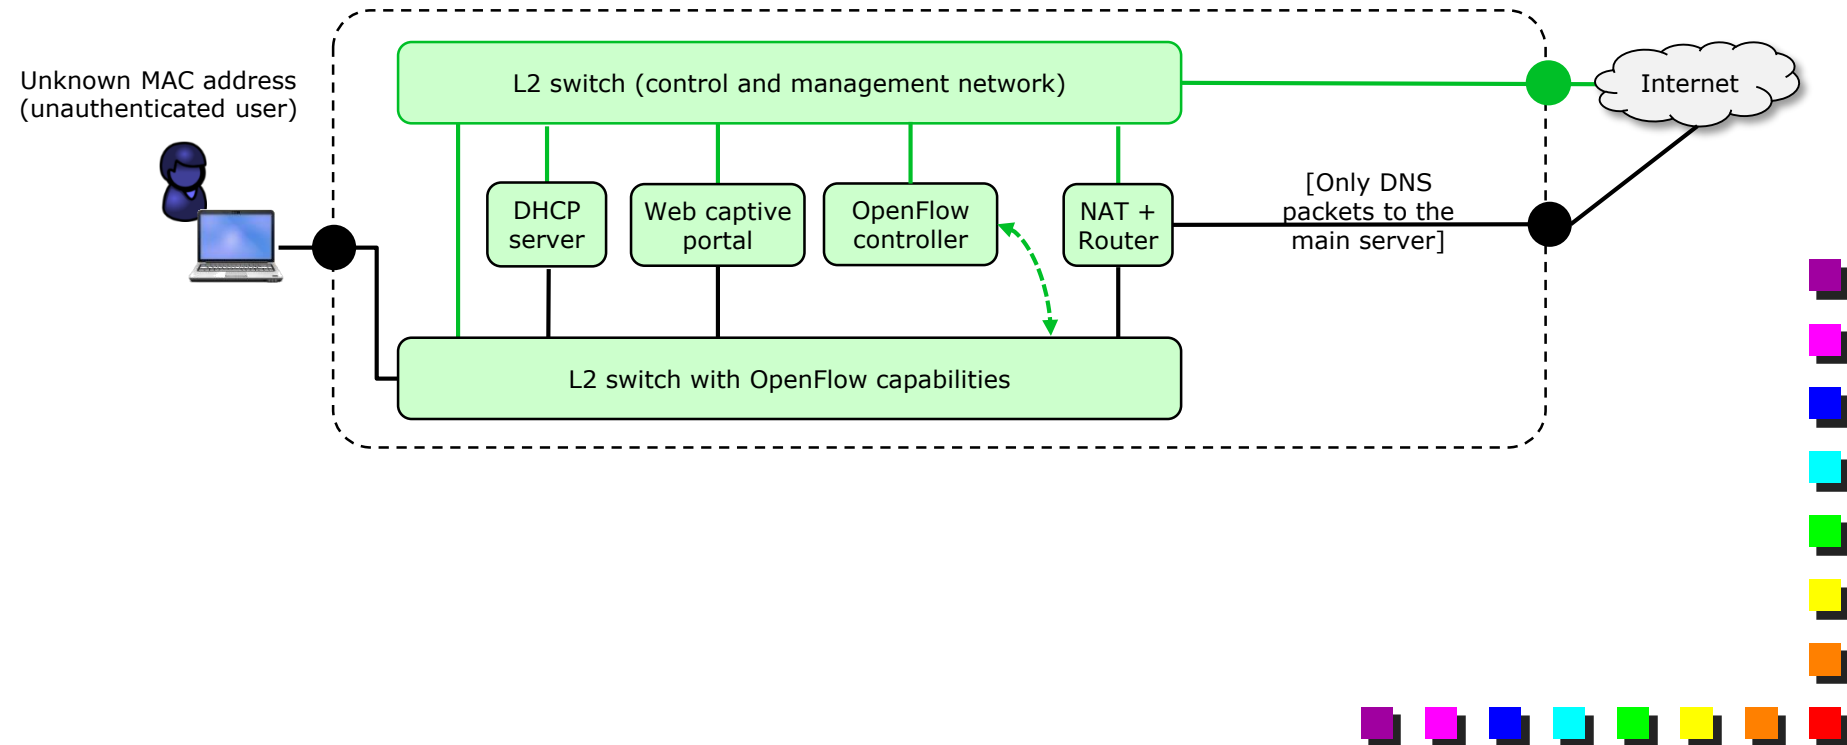
- Each **end user** is associated to a service graph operating on his own traffic
 - I.e., packets coming from / sent to his **MAC** address
 - The **network provider** forces the traffic of each user to cross an additional service graph under his control
 - Provides basic network services needed to connect to the Internet
 - When the user connects to the network:
 - is authenticated using a particular graph (**Auth_NF-FG**)
 - his own service graph **User_NF-FG** is instantiated on the Universal Node
 - His traffic is forced to cross the **User_NF-FG** first, then the **ISP_NF-FG** before reaching the Internet
- 
- A vertical column of colored squares (purple, magenta, blue, cyan, green, yellow, orange, red) on the right side of the slide.

Possible physical setup



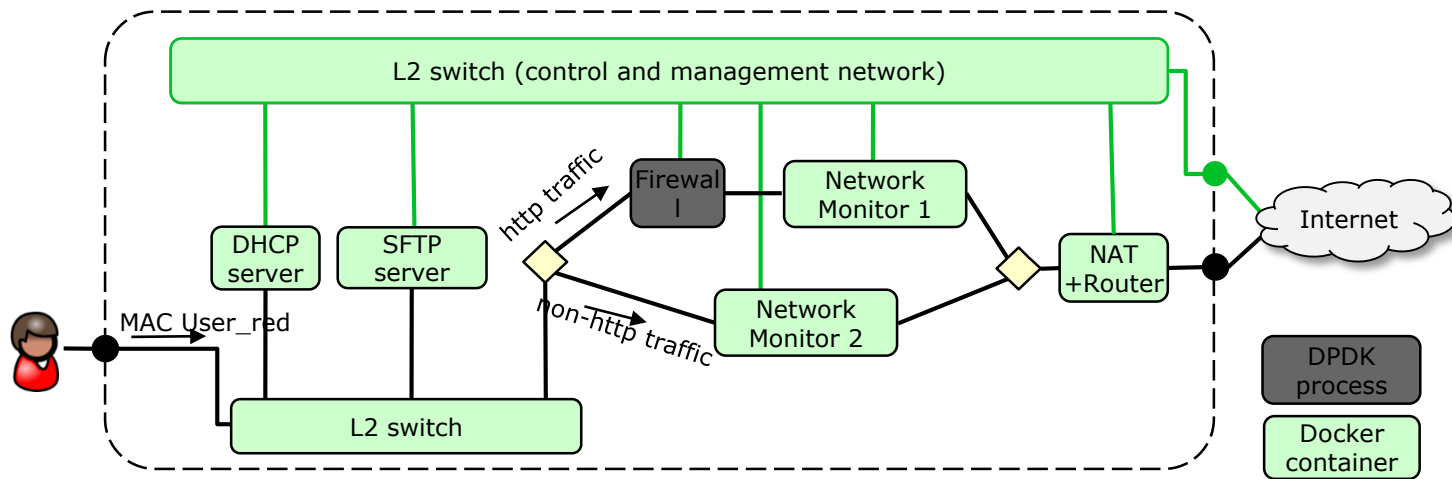
Authentication NF-FG (Auth_NF-FG)

- Handles the traffic generated by unknown devices (unknown MAC addresses)
- Provides a way to authenticate users



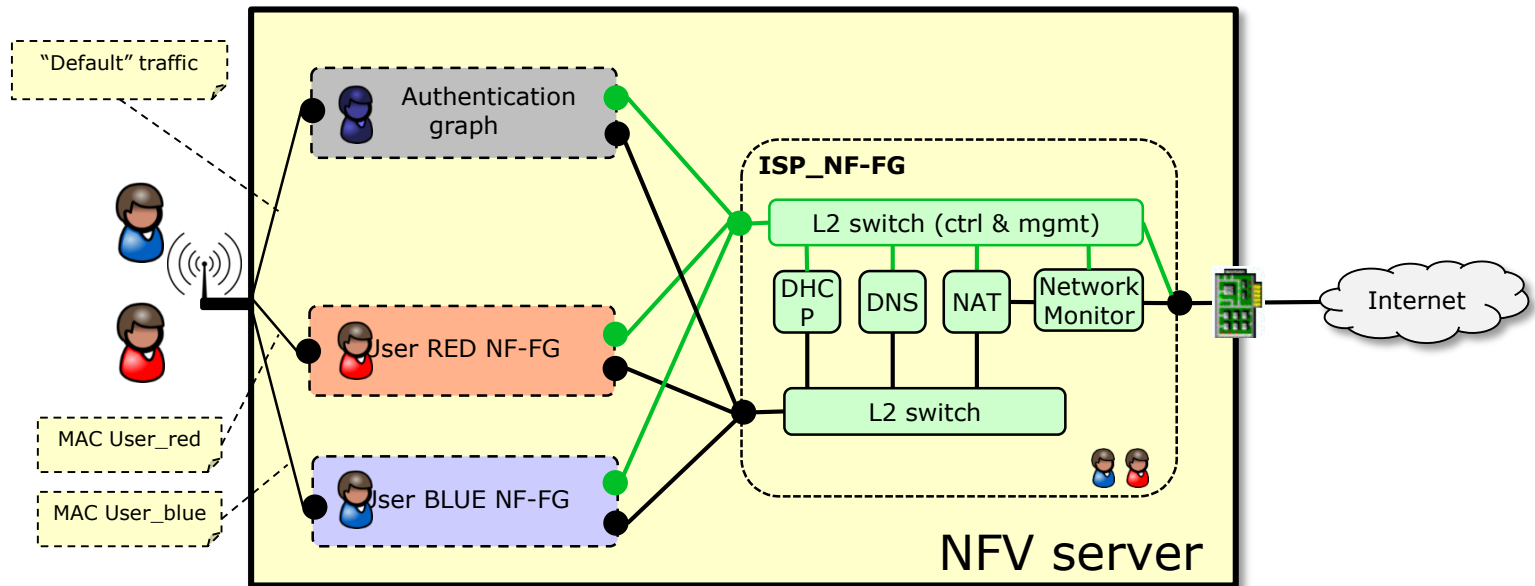
User_NF-FG

- Each user is associated to a specific NF-FG
 - Includes both transparent (e.g., firewall) and non-transparent (e.g., SFTP server) functions



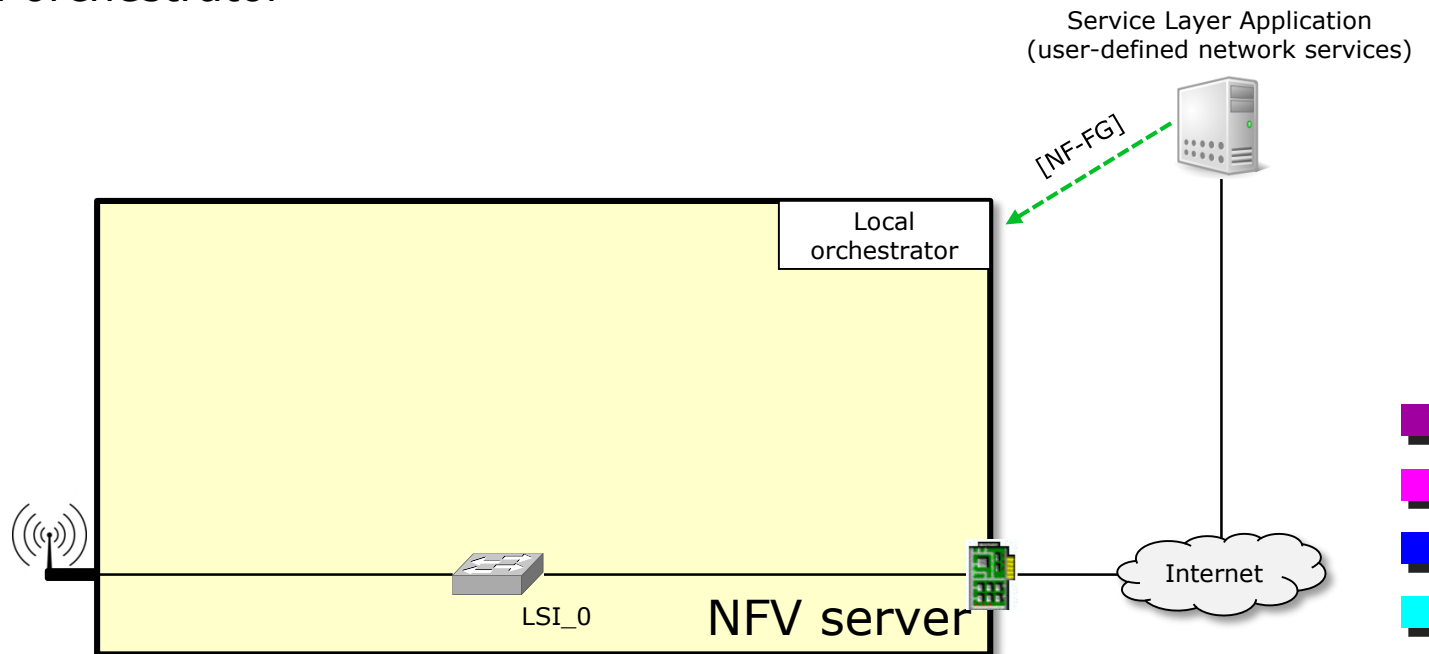
ISP_NF-FG

- Example of a possible set of Network Functions under the control of the ISP



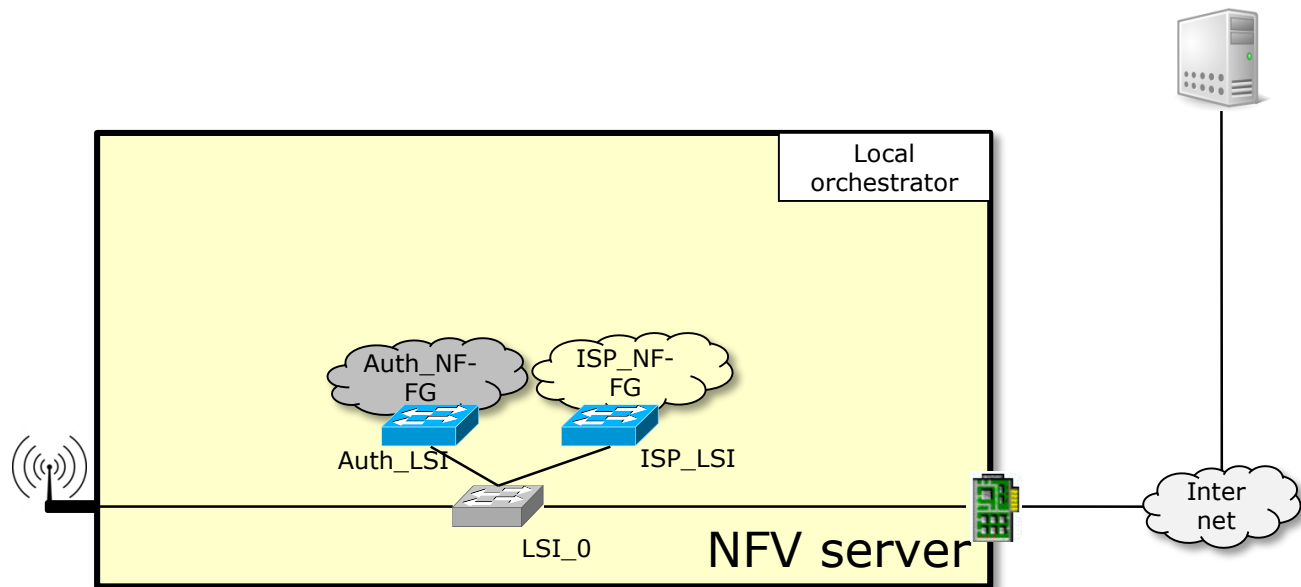
Example step 0: system startup

- Basic software running
 - Softswitch running (LSI0)
 - Local orchestrator



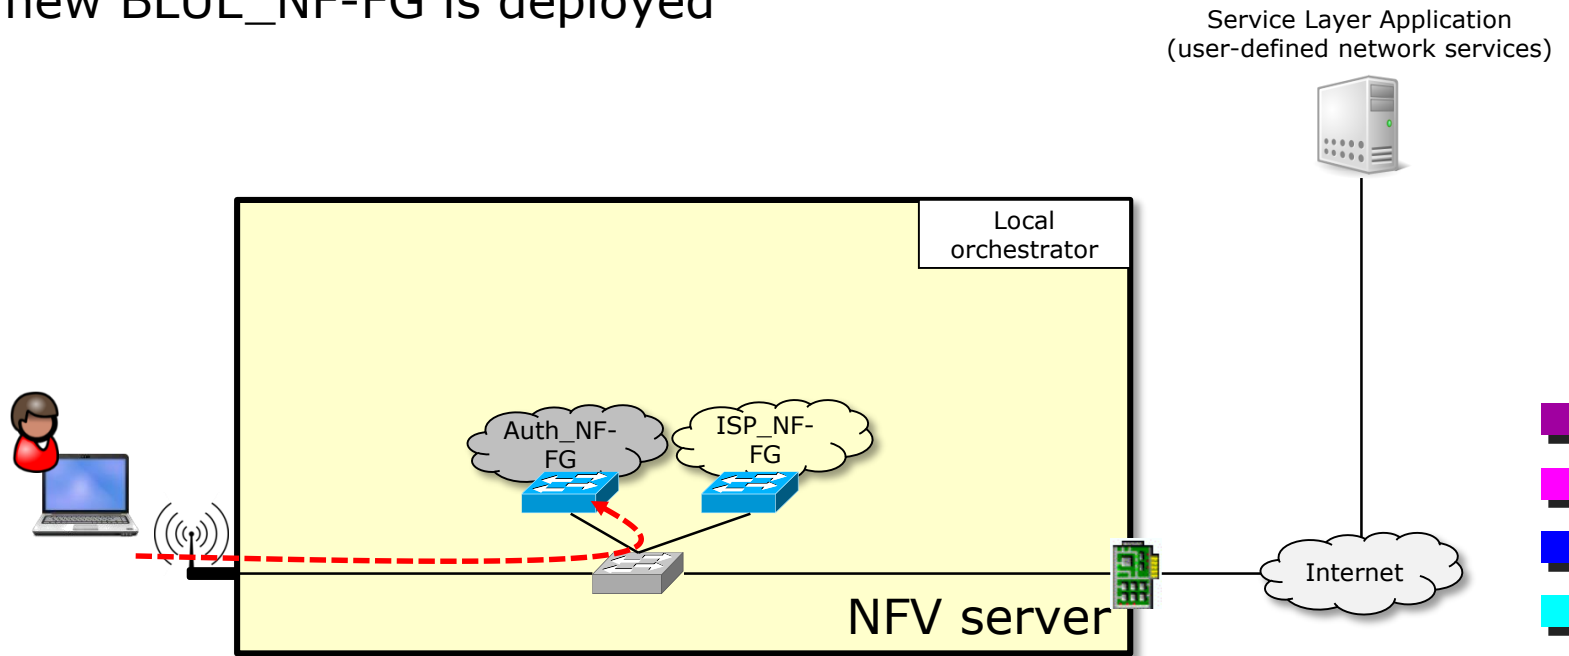
Example step 1: graph startup

- Auth_NF-FG and ISP_NF-FG automatically deployed
- Three Logical Switching Instances (LSI) active
- All incoming (from the WiFi) traffic sent to AuthLSI

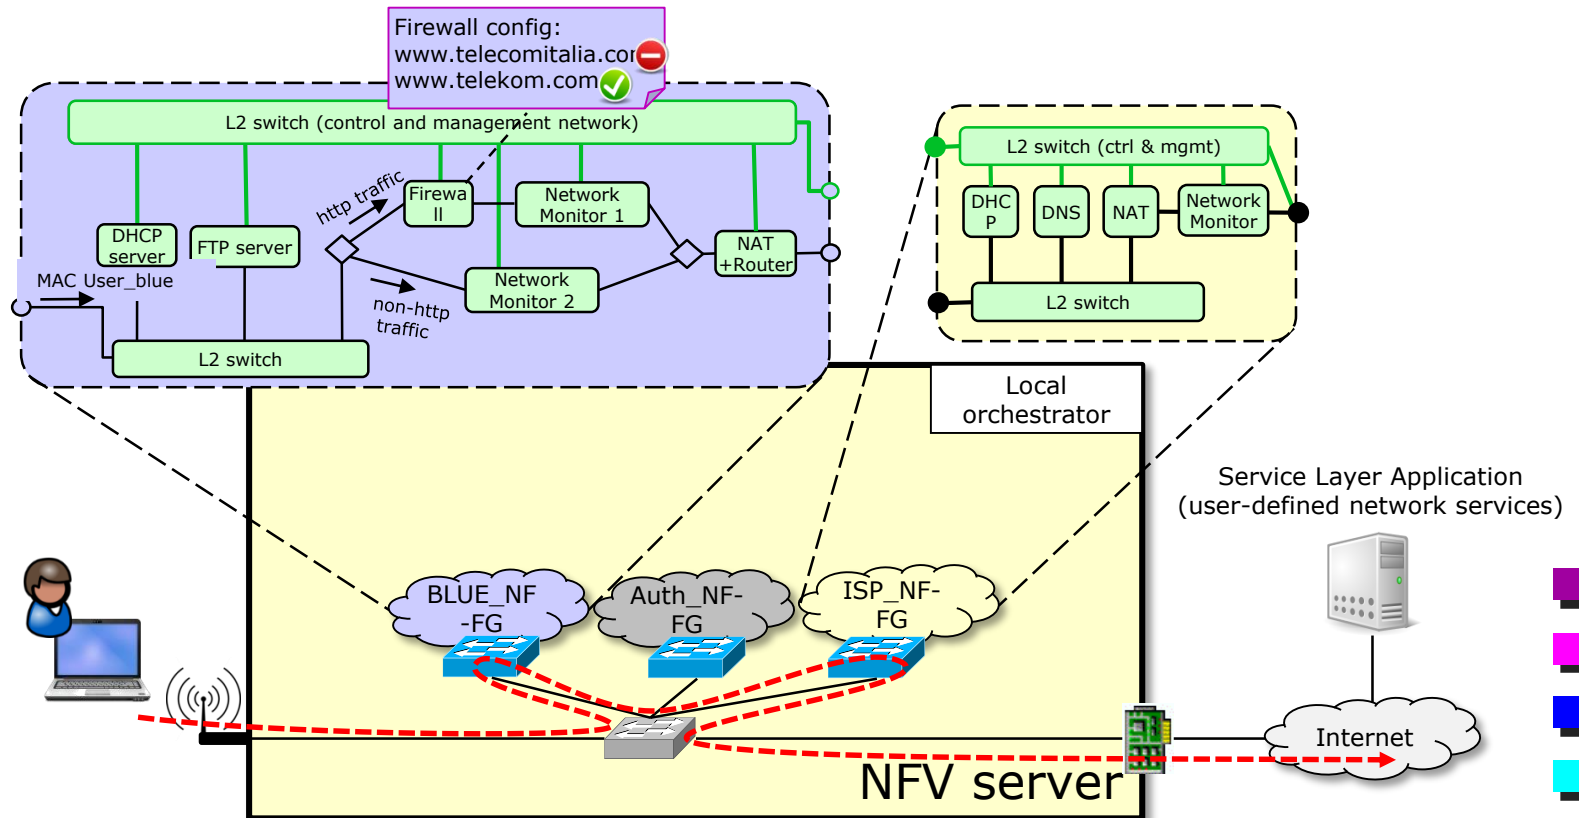


Example step 2: user (BLUE) connects

- User web traffic redirected to a captive portal
- User (BLUE) authenticates
- A new BLUE_NF-FG is deployed

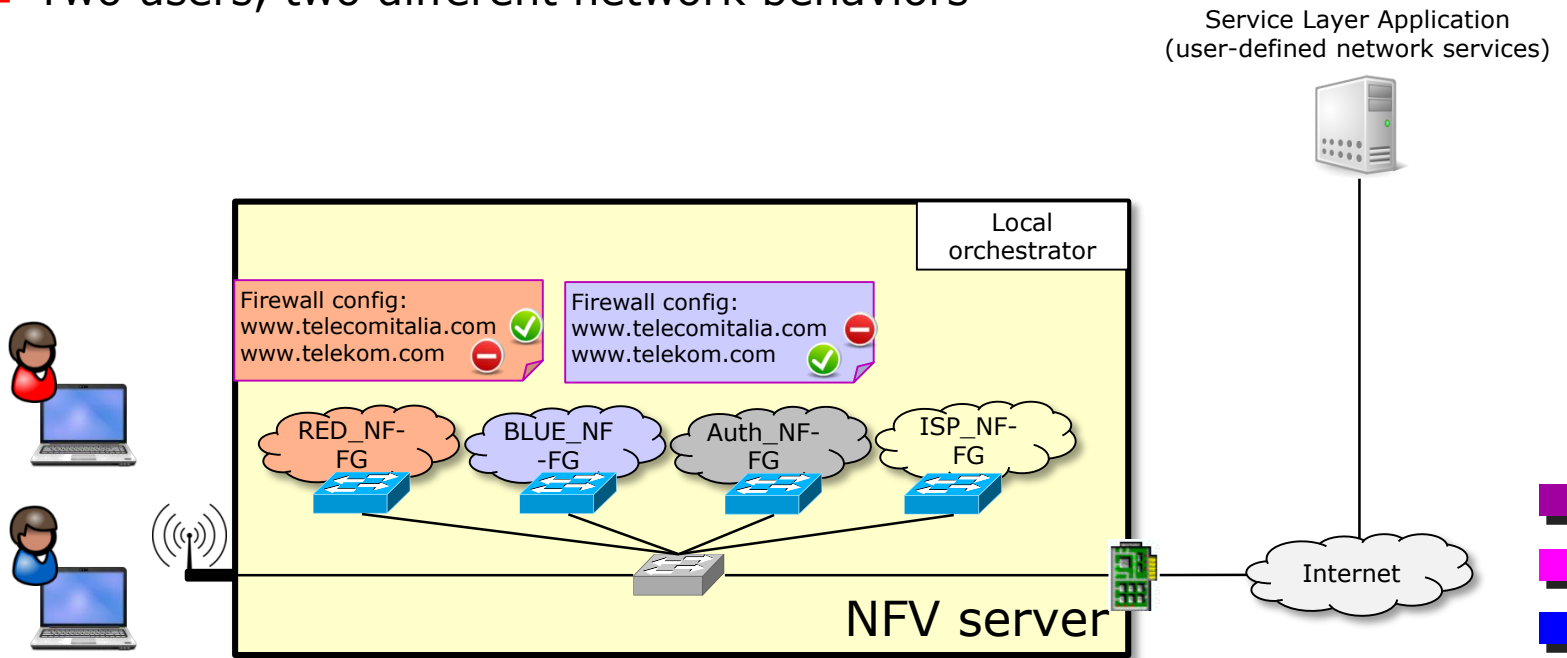


Example step 2b: user BLUE connected



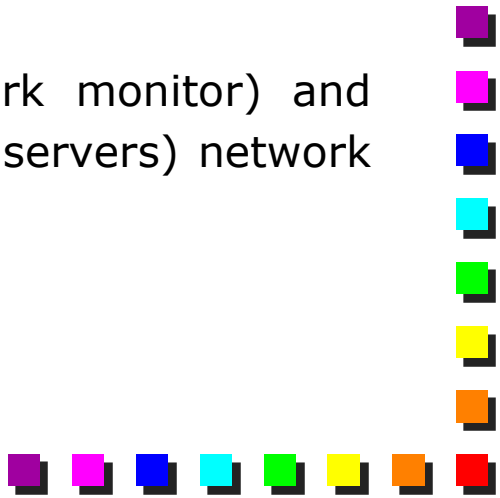
Example step 3: users BLUE and RED

- Logging in User RED from another laptop
 - Two users, two different network behaviors





Conclusions


- Shown the capability of an NFV server to deploy arbitrary network services, starting with a minimal set of components (softswitch, local orchestrator, NF-FG)
 - On-demand deployment of (user-defined) service graphs and run-time modifications of the attaching points
 - Support for a large number of network functions
 - Support for complex and cascading service graphs
 - Support for network functions with different architectures
 - Docker containers, DPDK native processes
 - Support for transparent (e.g., firewall, network monitor) and non-transparent (e.g., SFTP server, DHCP/DNS servers) network functions
- 



Additional content



Future integrated routers for NFV (1)



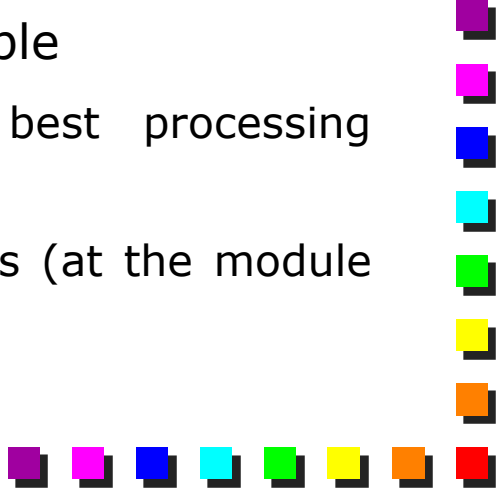
General
purpose
processing
linecards



Network
processing
linecards

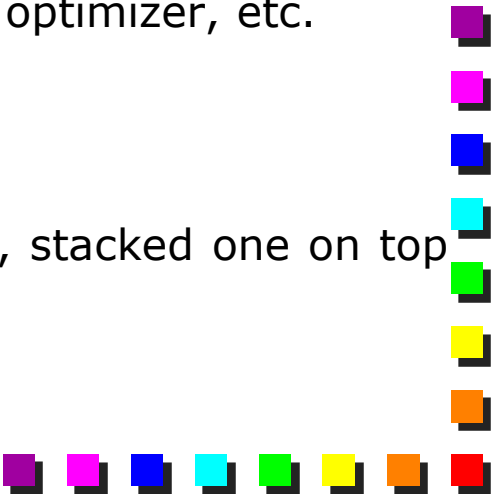



Future integrated routers for NFV (2)

- Routers with general purpose processing + network processing
 - More efficient (possibility to exploit hardware accelerators)
 - More scalable
 - Reduced network traffic to/from the router
 - Routers may become “hybrid” platforms supporting different kind of applications
 - From network applications to traditional VMs
 - Different hardware accelerators may be available
 - Each application can be mapped on the best processing component
 - Mapping can be done even at finer granularities (at the module level)
- 


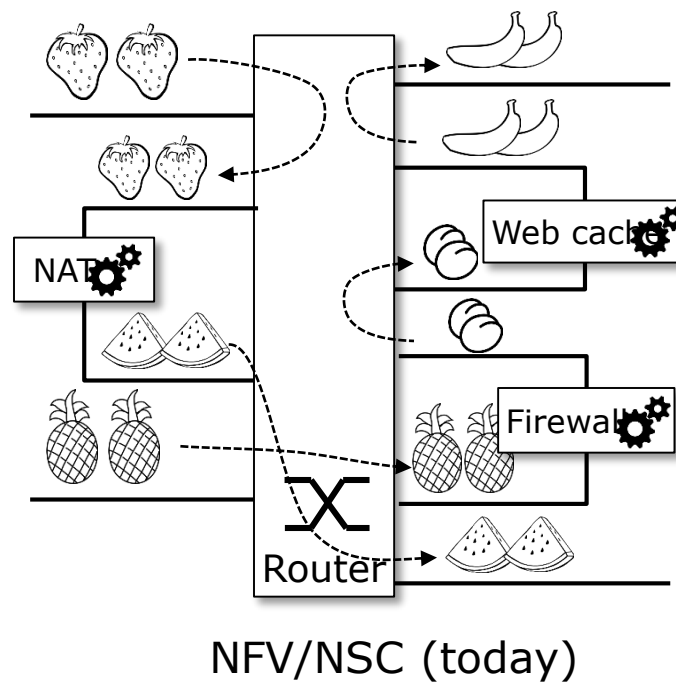



Some integrated routers available right now

- Router + x86 processing board tied in the same chassis
 - Interconnected through an high speed network (GbE)
 - Not too much integration
 - Possibility to register events in the data path that trigger an action on the x86 board
 - It reminds the old days of the Cat5K with switching and routing in the same chassis
 - Applications
 - Not very “network related”
 - Turn lights on/off, handle alarms, fax server, WAN optimizer, etc.
 - Not very successful
 - Not what we need
 - It was not so different from a router and a server, stacked one on top of the other
- 



SDN vs. NFV/SFC





Slim or fat routers?

- SDN predicts routers are slim
 - Slim routers suggest that the network will become a dumb pipe
 - Slim routers are good for “low value, low margins” vendors
 - Data plane performance requirements (and the relative simplicity of data plane tasks) suggest that (at least edge) routers need to be fat
 - Fat routers may suggest that the network will still have advanced processing capabilities
 - Hence, intelligence
 - Hence, value
 - Fat routers are good for current network vendors
- 